

# Parallel Programming

## With MATLAB Examples

Marek Rychlik

Department of Mathematics  
University of Arizona

February 18, 2023

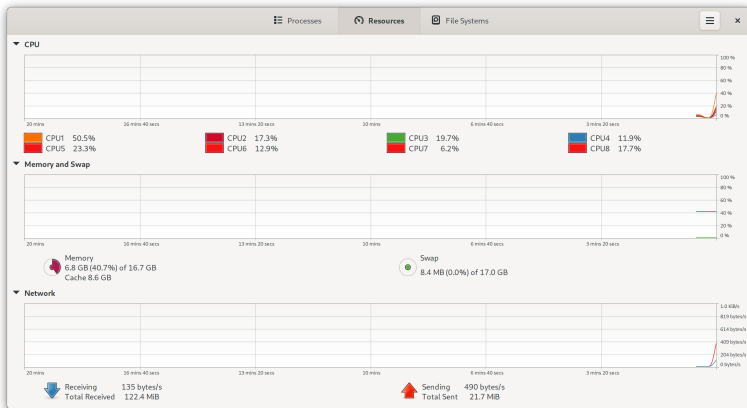
# Parallelism in the OS

- A modern OS has a multitude of processes running, as shown by a **system monitor**
- OS creates **an illusion of parallelism** even if it runs on a single CPU not capable of multi-threading in hardware.

Process Name	Status	Virtual Mem	Resident	Shared Mem	% CPU	CPU Time	Nice	ID	Memory	Waiting Channel	Control
systemd	Sleeping	23.9 MB	15.3 MB	10.8 MB	0.00	0:01.00	0	3865	4.5 MB	do_epoll_wait	user.slice
deja-dup	Running	790.1 MB	50.1 MB	36.4 MB	7.86	30:25.62	0	6051	13.8 MB	0	user.slice
duplcity	Running	534.5 MB	69.7 MB	18.7 MB	12.88	9:57.79	0	44244	50.9 MB	0	user.slice
gnome-shell	Running	5.9 GB	346.5 MB	170.2 MB	1.97	8:39.77	0	4031	176.3 MB	0	user.slice
gnome-system-monitor	Running	850.8 MB	115.8 MB	61.8 MB	8.52	0:19.66	0	45069	54.0 MB	0	user.slice
firefox	Running	4.2 GB	528.7 MB	227.5 MB	0.44	5:40.25	0	36150	301.1 MB	do_sys_poll	user.slice
Isolated Web Co	Running	3.5 GB	575.7 MB	125.3 MB	0.22	6:20.96	0	36593	450.4 MB	do_sys_poll	user.slice
Web Content	Sleeping	2.7 GB	92.1 MB	75.5 MB	0.00	0:00.19	0	44810	16.6 MB	do_sys_poll	user.slice
Web Content	Sleeping	2.7 GB	79.0 MB	66.6 MB	0.00	0:00.12	0	45045	12.4 MB	do_sys_poll	user.slice
Socket Process	Sleeping	250.7 MB	42.3 MB	33.5 MB	0.00	0:00.09	0	36261	8.8 MB	do_sys_poll	user.slice
Privileged Cont	Sleeping	2.8 GB	127.4 MB	97.1 MB	0.00	0:02.61	0	36283	30.3 MB	do_sys_poll	user.slice
RDD Process	Sleeping	331.6 MB	43.6 MB	33.1 MB	0.00	0:00.14	0	36474	10.5 MB	do_sys_poll	user.slice
Utility Process	Sleeping	333.2 MB	46.7 MB	35.3 MB	0.00	0:00.15	0	36475	11.4 MB	do_sys_poll	user.slice
Web Content	Sleeping	2.7 GB	78.8 MB	66.3 MB	0.00	0:00.11	0	44880	12.4 MB	do_sys_poll	user.slice
Web Content	Sleeping	2.7 GB	75.8 MB	63.8 MB	0.00	0:00.10	0	45172	12.1 MB	do_sys_poll	user.slice
WebExtensions	Sleeping	3.2 GB	299.9 MB	99.1 MB	0.00	0:29.70	0	36345	200.8 MB	do_sys_poll	user.slice
Xwayland	Sleeping	463.8 MB	117.7 MB	70.3 MB	0.00	1:44.88	0	4858	47.4 MB	do_epoll_wait	user.slice
emacs	Sleeping	1.0 GB	219.3 MB	98.0 MB	0.00	1:15.76	0	15196	121.3 MB	do_select	user.slice
MATLAB	Sleeping	11.3 GB	2.9 GB	470.3 MB	0.00	15:42.31	0	16966	2.5 GB	futex_wait_queue	user.slice
mwdocsearch	Sleeping	737.1 MB	15.1 MB	10.8 MB	0.00	0:03.35	0	17191	4.3 MB	do_epoll_wait	user.slice
gnome-screenshot	Sleeping	776.3 MB	48.7 MB	36.6 MB	0.00	0:00.61	0	45196	12.1 MB	do_sys_poll	user.slice
iconf-service	Sleeping	163.0 MB	9.2 MB	5.2 MB	0.00	0:00.27	0	4162	4.0 MB	do_sys_poll	user.slice
goc-identity-service	Sleeping	544.3 MB	9.0 MB	7.8 MB	0.00	0:04.16	0	4254	1.2 MB	do_sys_poll	user.slice
MATLABConnector	Sleeping	3.8 GB	161.7 MB	82.3 MB	0.00	10:22.73	0	4566	79.6 MB	futex_wait_queue	user.slice
MathWorksServiceHost	Sleeping	2.2 GB	107.8 MB	84.2 MB	0.00	1:16.64	0	4672	23.6 MB	futex_wait_queue	user.slice
gsd-housekeeping	Sleeping	539.1 MB	8.4 MB	6.7 MB	0.00	0:03.87	0	4325	1.7 MB	do_sys_poll	user.slice
xdg-desktop-portal-gtk	Sleeping	608.5 MB	30.5 MB	23.7 MB	0.00	0:01.27	0	4886	6.8 MB	do_sys_poll	user.slice

Figure: Explanations

# How many CPUs/Hardware threads do I have?



# Forking in Bash (&) — a minimal variant

```
1 (sleep 1e-2; echo -n "Hello , ") & \  
2   (sleep 1e-2; echo -n "World!")
```

# Forking in Bash (& I

OS

Parallelism

SIMD on GPU

Trivial

Parallelism

An Intro to

MPI

Troop

Counting

Example

Line graph topology

Tree graph topology

Implementation

Parallel

Mandelbrot

Wrapping up

```
1  #!/bin/bash
2  # EXAMPLE: Print 'Hello, ' and 'World!'
3  # in random order w/o a random number generator.
4  # HINT: We deliberately create a race condition.
5  if (($#)) ;then ntimes=$1 ;else ntimes=10; fi
6  function hello {
7      echo -n "Hello, "
8  }
9  function world {
10     echo -n "World!"
11 }
12 dlay=1e-2 # Change to 5 to see processes
13 for (( j=0; $j<$ntimes; j=$j+1 ))
14 do
15     # Fork with '&'
16     (sleep $dlay; hello) & (sleep $dlay; world)
17     echo " --Done with iteration: $j"
18 done
```

# Forking in Bash (&) II

OS

Parallelism

SIMD on GPU

Trivial

Parallelism

An Intro to  
MPITroop  
Counting  
ExampleLine graph topology  
Tree graph topology  
ImplementationParallel  
Mandelbrot

Wrapping up

```
[marek@cannonball]$ ./forkme.sh
Hello, World! --Done with iteration: 1
Hello, World! --Done with iteration: 2
World!Hello, --Done with iteration: 3
World!Hello, --Done with iteration: 4
Hello, World! --Done with iteration: 5
World!Hello, --Done with iteration: 6
Hello, World! --Done with iteration: 7
Hello, World! --Done with iteration: 8
Hello, World! --Done with iteration: 9
```

# A remarkable, more rare output

```
[marek@cannonball matlabmpi]$ ./forkme.sh
Hello, World! --Done with iteration: 1
Hello, World! --Done with iteration: 2
Hello, World! --Done with iteration: 3
Hello, World! --Done with iteration: 4
Hello, World! --Done with iteration: 5
Hello, World! --Done with iteration: 6
World! --Done with iteration: 7
Hello, World!Hello, --Done with iteration: 8
Hello, World! --Done with iteration: 9
```

# A Glossary of Terms I

**program counter** The location (address) of the instruction currently being executed; a place in a program

**process** A running program with all necessary resources (program counter, open file descriptors, memory state)

**fork, forking, clone** The UNIX/Linux **system call** which allows one process to create another one

**IPC, inter-process communication** The protocol by which two distinct processes can exchange information

**thread (of execution)** Formerly known as a **light-weight process** directly shares the state of memory (variables) with other threads; threads have **separate program counters**; a modern process is a **collection of threads**



# A Glossary of Terms II

**process/thread synchronization** Mechanisms by which one process tells another not to mess with some sensitive parts of its state; IPC can be used for process synchronization; threads are synchronized by **mutexes**

**mutex, futex** A mutually exclusive lock, which is a simple integer (logical) variable which is set/unset (=acquired/released) by a thread. What is important is the **interpretation** by another thread. A thread agrees not to do certain things when mutex is acquired by another, until it is released. **Semaphores** generalize mutexes to arbitrary integer values. **Futex** is a **fast mutex**, introduced by the Linux OS.

# A Glossary of Terms III

**atomicity** Some operations need to be atomic, such as changing the value of a mutex/semaphore. Atomicity means that a thread that reads the value of a mutex does not get an inconsistent value while another thread is **in the process of changing it**. Normal variables cannot be used as mutexes because reading and writing to them **is not atomic**. Atomicity is implemented using hardware (special instructions) and compiler (awareness that some variables must be changed atomically).

# Mandelbrot set definition

## Definition (The Mandelbrot set)

The set of complex numbers  $c$  for which the function  $f_c(z) = z^2 + c$  does not diverge to  $\infty$  when iterated from  $z = 0$ . In other words

$$\mathcal{M} = \left\{ c \in \mathbb{C} : \sup_n \left| \underbrace{f_c(f_c(\dots(0)\dots))}_{n \text{ times}} \right| < \infty \right\}$$

# The Mandelbrot set on GPU

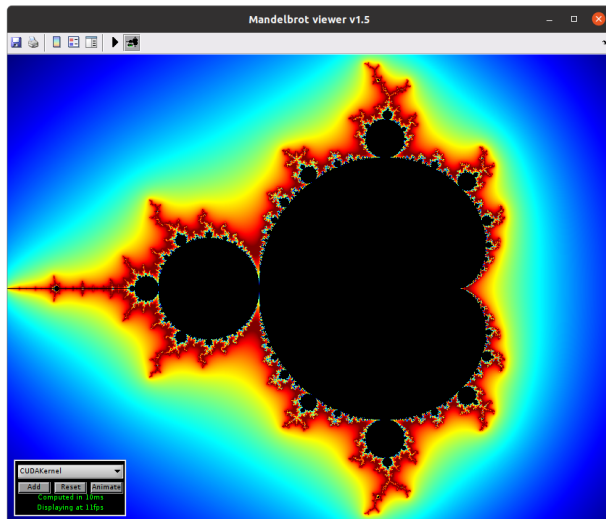


Figure: Example page: **Mandelbrot Set**

# The Julia definition

## Definition (The Julia set)

For fixed  $c$ , the set of complex numbers  $z$  for which the function  $f_c(z) = z^2 + c$  does not diverge to  $\infty$  when iterated from  $z$ . In other words

$$\mathcal{J}_c = \left\{ z \in \mathbb{C} : \sup_n \left| \underbrace{f_c(f_c(\dots(z)\dots))}_{n \text{ times}} \right| < \infty \right\}$$

# The Julia set on GPU




Figure: Example page: Look for a demo on MATLAB [File Exchange](#)

# MATLAB 'parfor' (parallel for) I

OS

Parallelism

SIMD on GPU

Trivial

Parallelism

An Intro to

MPI

Troop

Counting

Example

Line graph topology

Tree graph topology

Implementation

Parallel

Mandelbrot

Wrapping up

```

1 % FILE: parforEx.m
2 %p=gcp();
3 %mpilnit;
4 n=10;
5 % Evaluate x^2 for 1:n asynchronously and print results
6 parfor i=1:n
7     x=i^2
8     %pause(1);
9     disp([i,x]);
10 end
11 disp(' All done ');

```

parforEx

1	1
5	25
9	81
2	4
10	100
3	9
4	16
6	36
7	49

# MATLAB 'parfor' (parallel for) II

```
8      64
```

```
All done  
>>
```

## Question

Why does 'All done' print only once? Only at the end?



# Parallel pools, workers, clusters

OS

Parallelism

SIMD on GPU

Trivial

Parallelism

An Intro to  
MPI

Troop

Counting

Example

Line graph topology

Tree graph topology

Implementation

Parallel

Mandelbrot

Wrapping up

- TIP: first install **Parallel Computing Toolbox** and try its GUI to configure a cluster
- **Workers** are a MATLAB abstraction of threads, and they should directly map to hardware (CPU, hardware threads)
- A **parallel pool** is a collection of workers under the management of the **main thread**
- A parallel pool can live on one or more CPUs, and can be distributed across many computers; these details are abstracted away
- A **cluster** is defined by a configuration file (a **profile**, eg., 'local.settings') and it specifies computers and the number of CPU used on each machine. The configuration file must be placed in one of several standard places (see 'help parcluster').

# Accumulating values, reduction variables I

OS

Parallelism

SIMD on GPU

Trivial

Parallelism

An Intro to

MPI

Troop

Counting

Example

Line graph topology

Tree graph topology

Implementation

Parallel

Mandelbrot

Wrapping up

```

1 % FILE: reductionVar.m
2 % This file demonstrates a useful notion of a 'Reduction Variable'
3 % Makes it possible to accumulate values in a parfor without using
4 % spmd/gop.
5
6 p=gcp('nocreate');
7 if isempty(p)
8     p = parpool('local', 8)
9 end
10
11 disp(sprintf('Number of workers: %d', p.NumWorkers));
12
13 x=[];
14
15 parfor i = 1:10
16     pause(rand());
17     disp(i);
18     x = [x, i];
19 end
20
21 x =

>> reduction_var
Number of workers: 8

x =

```

# Accumulating values, reduction variables II

## Fact

Deterministic: the answer is always the same.

### Reduction Variables

R2018b

MATLAB® supports an important exception, called reduction, to the rule that loop iterations must be independent. A *reduction variable* accumulates a value that depends on all the iterations together, but is independent of the iteration order. MATLAB allows reduction variables in `par for`-loops.

Reduction variables appear on both sides of an assignment statement, such as any of the following, where `expr` is a MATLAB expression.

<code>X = X + expr</code>	<code>X = expr + X</code>
<code>X = X - expr</code>	See Associativity in Reduction Assignments in <a href="#">Requirements for Reduction Assignments</a>
<code>X = X .* expr</code>	<code>X = expr .* X</code>
<code>X = X * expr</code>	<code>X = expr * X</code>
<code>X = X &amp; expr</code>	<code>X = expr &amp; X</code>
<code>X = X   expr</code>	<code>X = expr   X</code>
<code>X = [X, expr]</code>	<code>X = [expr, X]</code>
<code>X = [X; expr]</code>	<code>X = [expr; X]</code>
<code>X = min(X, expr)</code>	<code>X = min(expr, X)</code>
<code>X = max(X, expr)</code>	<code>X = max(expr, X)</code>
<code>X = union(X, expr)</code>	<code>X = union(expr, X)</code>
<code>X = intersect(X, expr)</code>	<code>X = intersect(expr, X)</code>

# SPMD and SIMD

**SPMD** Stands for “Single program, multiple data”.

**Multiple autonomous processors**

simultaneously execute the same program at independent points (program counters). Can be implemented on general purpose CPUs (Intel, AMD)

**SIMD** Stands for “Single-instruction, multiple data”. A **vector processor** processes the same instruction on different data (example: coordinatewise addition or multiplication of two vectors).

Modern CPU(s) implements **both paradigms**:

- SIMD uses Intel/AMD **SSE instructions** and **vector registers**;
- SPMD uses multiple **threads**, **cores** and CPUs.

# The MPI (Message Passing Interface)

- The most successful realization of SPMD; used in MATLAB; 40 years of history
- Implementations in C, C++, Fortran exist, with high-level language interfaces (e.g., **Python**).
- Worker becomes a **lab**
- Worker knows its identity, or **labindex**
- The main thread is now a lab with **labindex==1** (recall: MATLAB has 1-based arrays)
- Labs communicate by using **collective communications**: labSend, labReceive, labSendReceive;
- synchronization: labBarrier, labBroadcast
- Labs can be organized as a graph with variable topology, e.g. edges of a hypercube, for the purpose of communicating with some neighbors

# Unintended blocking — a show stopper

- A **blocking operation** is one that stops the execution of the program (thread, process) until some condition is met
- Example: reading from a file. We wait for the data to be available (e.g., read from disk or network)
- Example: waiting for a mutex to be released
- **labReceive**, **labBarrier** are blocking operations
- A **non-blocking operation** does not wait for the condition to be met but immediately continues with the execution, reporting status to the caller
- Example: reading from a file in non-blockin mode reports the number of bytes successfully read. One repeatedly reads from the file, getting the file in chunks, until the end-of-file marker is found

# Deadlock (deadly embrace)

## Definition (Deadlock)

**Deadlock** (which is sometimes called **the deadly embrace**) occurs when two or more programs (threads, workers, labs) are each waiting for the others to complete a task before proceeding.

*The programs act like the overly congenial gophers in some Looney Tunes cartoons:*

*"Oh please, you first," says one.*

*"No no, I insist, you first," says the other.*

*And nothing goes anywhere.*

*Michael Meehan, Computerworld, Oct 29, 2001*

# An 'spmd' example (WRONG!) I

OS

Parallelism

SIMD on GPU

Trivial

Parallelism

An Intro to

MPI

Troop

Counting

Example

Line graph topology

Tree graph topology

Implementation

Parallel

Mandelbrot

Wrapping up

```

1  % FILE: race1.m
2  % This file demonstrates a simple race condition
3  % when trying to share a value between all workers
4  p=gcp('nocreate');
5  if isempty(p)
6      p = parpool('local', 8)
7  end
8
9  disp(sprintf('Number of workers: %d', p.NumWorkers));
10
11 value = Composite();
12
13 % An incorrect way to broadcast a value and
14 % receive it in all workers
15
16 spmd
17     pause(rand()./10);
18     if labindex == 1
19         for w=1:numlabs
20             display(sprintf('%d sending 7 to %d',labindex,w));
21             labSend(7,w);
22         end
23     end
24     value = labReceive;
25     display(sprintf('%d received %d from 1',labindex,value));
26 end
27
28 for w=1:p.NumWorkers
29     disp(value{w});
30 end

```



# An 'spmd' example (WRONG!) II

OS

Parallelism

SIMD on GPU

Trivial

Parallelism

**An Intro to  
MPI**

Troop

Counting

Example

Line graph topology

Tree graph topology

Implementation

Parallel

Mandelbrot

Wrapping up

```
>> race1
Number of workers: 8
Worker 1:
  1 sending 7 to 1
Error using race1
Error detected on worker 1.

Caused by:
  Error using race1
  Destination (1) is same as source, would cause deadlock.

>>
```

# An 'spmd' example (CORRECT!) I

```

1  % FILE: race1Fixed.m
2  % This file demonstrates a simple race condition
3  % when trying to share a value between all workers.
4  % NOTE: We avoid sending to ourselves. This
5  % avoids the race condition.
6
7  p=gcp('nocreate');
8  if isempty(p)
9      p = parpool('local', 8)
10 end
11
12 disp(sprintf('Number of workers: %d', p.NumWorkers));
13
14 value = Composite();
15
16 % A correct way to broadcast a value and
17 % receive it in all workers. However, lab 1 runs in O(n) time,
18 % so it is not an efficient way to broadcast data to others.
19
20 spmd
21     pause(rand()./10);
22
23     if labindex == 1
24         for w=2:numlabs
25             display(sprintf('%d sending 7 to %d',labindex,w));
26             labSend(7,w)
27         end
28         value = 7;

```

# An 'spmd' example (CORRECT!) II

```

29     else
30         value = labReceive;
31         display(sprintf('%d received %d from 1',labindex,value));
32     end
33
34 end
35
36 for w=1:p.NumWorkers
37     disp([w,value{w}]);
38 end

```

```

>>
Number of workers: 8
Worker 1:
    1 sending 7 to 2
    1 sending 7 to 3
    1 sending 7 to 4
    1 sending 7 to 5
    1 sending 7 to 6
    1 sending 7 to 7
    1 sending 7 to 8
Worker 2:
    2 received 7 from 1
Worker 3:
    3 received 7 from 1
Worker 4:
    4 received 7 from 1
Worker 5:

```

# An 'spmd' example (CORRECT!) III

```
5 received 7 from 1
Worker 6:
6 received 7 from 1
Worker 7:
7 received 7 from 1
Worker 8:
8 received 7 from 1
1      7
2      7
3      7
4      7
5      7
6      7
7      7
8      7
```

>>

## Broadcasting (another fix) I

Marek Rychlik

OS

Parallelism

SIMD on GPU

Trivial

Parallelism

An Intro to  
MPI

Troop

Counting

Example

Line graph topology

Tree graph topology

Implementation

Parallel

Mandelbrot

Wrapping up

```

1 % FILE: race1FixedBroadcast.m
2 % This file demonstrates a simple race condition
3 % when trying to share a value between all workers
4 % NOTE: In this version, we avoid the race condition
5 % by using labBroadcast.
6
7 p=gcp('nocreate');
8 if isempty(p)
9     p = parpool('local', 8)
10 end
11
12 disp(sprintf('Number of workers: %d', p.NumWorkers));
13
14 value = Composite();
15
16 % An incorrect way to broadcast a value and
17 % receive it in all workers
18
19 root=1;
20 spmd
21     pause(rand()./10);
22     if labindex == root
23         value = 7;
24         value = labBroadcast(root, value);
25         display(sprintf('Root==%d broadcast %d', labindex, value));
26     else
27         value = labBroadcast(root, 666); % Second input ignored on root
28         display(sprintf('%d received %d from root==%d', labindex, value, root));
29     end
30 end
31

```

## Broadcasting (another fix) II

Marek Rychlik

OS

Parallelism

SIMD on GPU

Trivial

Parallelism

An Intro to  
MPITroop  
Counting  
Example

Line graph topology

Tree graph topology

Implementation

Parallel  
Mandelbrot

Wrapping up

```

32 for w=1:p.NumWorkers
33     disp([w, value{w}]);
34 end

```

```

>>
Number of workers: 8
Worker 1:
    Root==1 broadcast 7
Worker 2:
    2 received 7 from root==1
Worker 3:
    3 received 7 from root==1
Worker 4:
    4 received 7 from root==1
Worker 5:
    5 received 7 from root==1
Worker 6:
    6 received 7 from root==1
Worker 7:
    7 received 7 from root==1
Worker 8:
    8 received 7 from root==1
    1      7

    2      7

    3      7

    4      7

    5      7

```

6      7


7      7

8      7

>>

# Broadcasting (another fix) III

# Linear Troop Topology



**Figure:** A line of soldiers counting themselves using message-passing rule-set A. The commander can add “3” from the soldier in front, “1” from the soldier behind, and “1” for himself, and deduce that there are 5 soldiers in total.



# Message-passing rule-set A (parallel pseudo-code).

- 1 If you are the front soldier in the line, say the number **one** to the soldier behind you.
- 2 If you are the rearmost soldier in the line, say the number **one** to the soldier in front of you.
- 3 If a soldier ahead of or behind you says a number to you, add one to it, and say the new number to the soldier on the other side.

## Implementation I

OS

Parallelism

SIMD on GPU

Trivial

Parallelism

An Intro to

MPI

Troop

Counting

Example

Line graph topology

Tree graph topology

Implementation

Parallel

Mandelbrot

Wrapping up

```

1  % FILE:      soldiers.m
2  % Mackay, algorithm 16.1: to count soldiers
3  % marching in line
4  %
5  % NOTE: If you are running the program on a
6  % processor with n=4 cores and 2*n
7  % hyperthreads, the setting for the 'local'
8  % cluster is used, and the number of workers
9  % in a parpool is set automatically to n,
10 % ignoring hyperthreading. You can modify
11 % the number of worker threads using Matlab
12 % GUI, using Home > Parallel > Manage Cluster
13 % Profiles > Edit. So, if you request 8
14 % workers, make sure to first edit the local
15 % profile and increase the number of allowed
16 % workers to >=8. I changed it to 64.
17 %
18 p = gcp('nocreate');
19 numSoldiers=5;
20 % Must have at least numSoldiers workers
21 if ~isempty(p) && p.NumWorkers < numSoldiers
22     delete(p);
23     p=[];
24 end
25 if isempty(p)
26     % Create a local parpool with num. workers == num. soldiers
27     p = parpool('local',numSoldiers);
28 end
29 mpiInit;
30 commander=2;
31 % Must have at least commander+1 workers

```

## Implementation II

OS

Parallelism

SIMD on GPU

Trivial

Parallelism

An Intro to

MPI

Troop

Counting

Example

Line graph topology

Tree graph topology

Implementation

Parallel

Mandelbrot

Wrapping up

```

32 assert(p.NumWorkers > commander);
33 spmd
34     me=labindex;
35     value=0;
36     if me==commander
37         [value1,source1,tag1]=labReceive;
38         fprintf( '%d=commander got %d from %d\n',me,value1,source1);
39         [value2,source2,tag2]=labReceive;
40         fprintf( '%d=commander got %d from %d\n',me,value2,source2);
41         value=value1+value2+1;
42         fprintf( '%d=commander says: count is %d\n',me,value);
43     elseif me==1
44         value=1;
45         dest=2;
46         fprintf( '%d sending %d to %d\n',me,value,dest);
47         labSend(value,dest);
48     elseif me==numlabs
49         value=1;
50         dest=me-1;
51         fprintf( '%d sending %d to %d\n',me,value,dest);
52         labSend(value,dest);
53     else
54         [value,source,~]=labReceive;
55         value=value+1;
56         if source==me-1
57             dest=me+1;
58         elseif source==me+1
59             dest=me-1;
60         end
61         fprintf( '%d sending %d to %d\n',me,value,dest);
62         labSend(value,dest);

```

# Implementation III

# General topology of the troop




Figure: A swarm of guerillas.

# Arranging workers/labs into a graph I

```

1  % FILE:    buildTroop.m
2
3  % Efficient graph encoding (often used in MPI programs)
4  % Adjacency matrix; passed as 1-d array, in which each vertex is followed
5  % Row format: node followed by neighbors, sentinel 0.
6  % A final zero is added to terminate the structure.
7  adj=[1,2,0,...
8        2,1,3,11,0,...
9        3,2,4,5,0,...
10       4,3,0,...
11       5,3,0,...
12       6,8,0,...
13       7,8,0,...
14       8,6,7,9,0,...
15       9,8,10,12,0,...
16       10,9,11,0,...
17       11,2,10,0,...
18       12,9,13,14,0,...
19       13,12,0,...
20       14,12,0,...
21       0];
22
23 % Automatically determine te number of soldiers (nodes)
24 numSoldiers = numel(find(adj==0))-1;
25 commander = 9; % Designate the commander
26
27 % Start the parpool (thread pool)
28 p = gcp('nocreate');
```

# Arranging workers/labs into a graph II

```

29 if ~isempty(p) && p.NumWorkers ~= numSoldiers
30     delete(p);
31     p=[];
32 end
33 if isempty(p)
34     p = parpool('local',numSoldiers);
35 end
36
37 mpiinit;
38
39 % Convert nb to cell array
40 nb=Composite();
41 start=1;
42 for s=1:numSoldiers;
43     me=adj(start);
44     neighbors=[];
45     n=start+1;
46     % Make a list of neighbors
47     while adj(n)~=0
48         neighbors=[neighbors, adj(n)];
49         n=n+1;
50     end
51     nb{me}=neighbors;
52     start=n+1;
53 end
54 assert(adj(start)==0);
55
56


```

# Arranging workers/labs into a graph III

```
57 % A normal function call to let soldiers report the neighbors
58 report(nb);
59
60 % A consistency check for graph data
61 l=adjacency_matrix(nb);
62 % Check symmetry of the adjacency relation
63 assert(all(l==l')));
64 % Check for 'no loops' (loop=connection of edge to itself)
65 assert(all(diag(l)==0));
66 % Check for 'no cycles'; upper triangular portion of l should be nilpotent
67 assert(all(all(triu(l)^numSoldiers==0)));
68
69 %g = graph(l);
70 %plot(g, 'LineWidth',4, 'NodeFontSize',44, 'MarkerSize',5, 'NodeLabelColor', 'blue', '
    NodeFontWeight', 'bold');
```



# A tree topology of the troop



# Message-passing rule-set B. I

OS

Parallelism

SIMD on GPU

Trivial

Parallelism

An Intro to  
MPITroop  
Counting  
ExampleLine graph topology  
Tree graph topology  
ImplementationParallel  
Mandelbrot

Wrapping up

```

1: procedure MESSAGEPASSING(Graph)
2:    $N \leftarrow$  the count your neighbours in Graph
3:    $m \leftarrow 0$  ▷ count of messages received from
      neighbours
4:   for  $j$  from 1 to  $N$  do
5:      $v_j \leftarrow -1$  ▷ initial value of message is invalid
6:   end for
7:    $V \leftarrow 0$  ▷ running total of messages you have
      received
8:   if  $m == N - 1$  then
9:     Find neighbor  $j$  such that  $v_j == -1$  ▷ the only
      one who has not sent you a message
10:    Tell them the number  $V + 1$ 
11:  end if

```

## Message-passing rule-set B. II

OS

Parallelism

SIMD on GPU

Trivial

Parallelism

An Intro to  
MPITroop  
Counting  
ExampleLine graph topology  
Tree graph topology  
ImplementationParallel  
Mandelbrot

Wrapping up

```
12:   if  $V == N$  then
13:       the number  $V + 1$  is the required total.
14:       for each neighbour  $n$  do
15:           say to neighbour  $n$  the number  $V + 1 - v_n$ 
16:       end for
17:   end if
18: end procedure
```

## Implementation I

```

1  % FILE: soldiers2.m
2  % NOTE: Make sure to modify the parpool to allow 14 workers.
3
4  buildTroop;
5
6  % Main course: count the soldiers by message passing
7  spmd
8      me=labindex;
9      N=length(nb);           % Neighbor count
10     m=0;                    % Message count
11     v=ones(N,1);            % Message values
12     V=0;                    % Running total of messages
13
14     labBarrier;              % Not needed, harmless, for demo purposes
15
16     % Receive first N-2 messages
17     while m < N-1
18         [isDataAvail, source]=labProbe;
19         if isDataAvail       % If available, get the data
20             n=find(source==nb,1); % Find which neighbor sent the message
21             assert(~isempty(n)); % Otherwise it is not a neighbor
22             fprintf('%d sees data available from %d...\n ',me,source);
23             value=labReceive(source);
24             fprintf('%d received value %d from %d.\n ',me,value,source);
25             m=m+1;
26             v(n)=value;
27             V=V+value;
28         end
29     end
30
31     %labBarrier;              % Will break the code!!!

```

## Implementation II

```

32 assert(m==N-1); % Check number of messages
33
34 % Send the message to who has not send us a message
35 n=find(v==-1,1); % Identify who has not send us a message
36 dest=nb(n);
37 value_to_send=V+1;
38 fprintf(' %d sending %d to %d...\n',me,value_to_send,dest);
39 labSend(value_to_send,dest);
40 %labSendReceive(value_to_send,dest);
41 fprintf(' %d completed sending %d to %d.\n',me,value_to_send,dest);
42 fprintf(' %d waiting for message from %d...\n',me,dest);
43 [value,last_source]=labReceive(dest);
44 fprintf(' %d received %d from %d.\n',me,value,last_source);
45 assert(last_source==dest); % Last message source
46 v(n)=value;
47 V=V+value;
48 m=m+1;
49
50 %labBarrier; % Will break the code!!!
51
52 assert(m==N); % Check message count
53
54 % Send message to everyone except the one who was the last to send us
55 % a message.
56 for l=1:N
57     if l==n
58         continue; % Do not send again
59     end
60     value_to_send=V+1-v(l);
61     fprintf(' %d sending %d to %d...\n',me,value_to_send,nb(l));
62

```

## Implementation III

OS

Parallelism

SIMD on GPU

Trivial

Parallelism

An Intro to  
MPI

Troop

Counting

Example

Line graph topology

Tree graph topology

Implementation

Parallel

Mandelbrot

Wrapping up

```
63         labSend(value_to_send,nb(l));
64         fprintf( '%d completed sending %d to %d.\n',me,value_to_send,nb(l));
65     end
66     if me==commander
67         fprintf( 'COMMANDER %d reporting count of %d.\n', me, V+1);
68     end
69     fprintf( '%d is done.\n',me);
70 end
71
72 % A demonstration that a Composite is a kind of cell array
73 % Print the totals of all soldiers.
74 for n=1:numSoldiers
75     fprintf( 'Running total of %d is %d.\n', n, V{n});
76 end
```

## A modified implementation I

```

OS
Parallelism
SIMD on GPU
Trivial
Parallelism
An Intro to
MPI
Troop
Counting
Example
Line graph topology
Tree graph topology
Implementation
Parallel
Mandelbrot
Wrapping up

1  % FILE: soldiers3.m
2  % This is like soldiers2, but uses labSendReceive instead
3  % of 2 calls labSend/labReceive. It is preferred to do it this way,
4  % as there is a smaller chance of programming error causing a race
5  % condition (labReceive when there is noone sending, or labSend when
6  % there is noone waiting to labReceive).
7
8  buildTroop;
9
10 % Main course: count the soldiers by message passing
11 spmd
12     me=labindex;
13     N=length(nb); % Neighbor count
14     m=0; % Message count
15     v=-ones(N,1); % Message values
16     V=0; % Running total of messages
17
18     fprintf(' %d reached barrier.\n', me);
19     labBarrier; % Not needed, harmless, for demo purposes
20     fprintf(' %d crossed barrier.\n', me);
21
22 % Receive first N-2 messages
23 while m < N-1
24     [isDataAvail, source]=labProbe;
25     if isDataAvail % If available, get the data
26         n=find(source==nb,1); % Find which neighbor sent the message
27         assert(~isempty(n)); % Otherwise it is not a neighbor
28         fprintf(' %d sees data available from %d...\n ',me,source);
29         value=labReceive(source);
30         fprintf(' %d received value %d from %d.\n ',me,value,source);
31         m=m+1;

```

## A modified implementation II

Marek Rychlik

OS 32  
 Parallelism 33  
 SIMD on GPU 34  
 35  
 Trivial 36  
 Parallelism 37  
 38  
 An Intro to 39  
 MPI 40  
 41  
 Troop 42  
 Counting 43  
 Example 44  
 Line graph topology 45  
 Tree graph topology 46  
 Implementation 47  
 48  
 Parallel 49  
 Mandelbrot 50  
 51  
 Wrapping up 52  
 53  
 54  
 55  
 56  
 57  
 58  
 59  
 60  
 61  
 62

```

32         v(n)=value;
33         V=V+value;
34     end
35 end
36
37 %labBarrier; % Will break the code!!!
38
39 assert(m==N-1); % Check number of messages
40
41 % Send the message to who has not send us a message
42 n=find(v==-1,1); % Identify who has not send us a message
43 dest=nb(n);
44 fprintf( '%d noticed not receiving from %d', dest);
45 value_to_send=V+1;
46 fprintf( '%d sending %d to %d...\n',me,value_to_send,dest);
47 value = labSendReceive(dest, value_to_send);
48 fprintf( '%d received %d from %d.\n',me,value,dest);
49 v(n)=value;
50 V=V+value;
51 m=m+1;
52
53 %labBarrier; % Will break the code!!!
54
55 assert(m==N); % Check message count
56
57 % Send message to everyone except the one who was the last to send us
58 % a message.
59 for l=1:N
60     if l==n
61         continue; % Do not send again
62     end

```



## A modified implementation III

OS

Parallelism

SIMD on GPU

Trivial

Parallelism

An Intro to

MPI

Troop

Counting

Example

Line graph topology

Tree graph topology

Implementation

Parallel

Mandelbrot

Wrapping up

```

63     value_to_send=V+1-v(l);
64     fprintf( '%d sending %d to %d...\n',me,value_to_send,nb(l));
65     labSend(value_to_send,nb(l));
66     fprintf( '%d completed sending %d to %d.\n',me,value_to_send,nb(l));
67 end
68 if me==commander
69     fprintf( 'COMMANDER %d reporting count of %d.\n', me, V+1);
70 end
71 fprintf( '%d is done.\n',me);
72 end
73
74
75 % A demonstration that a Composite is a kind of cell array
76 % Print the totals of all soldiers.
77 for n=1:numSoldiers
78     fprintf( 'Running total of %d is %d.\n', n, V{n});
79 end

```

# Sample output I

OS

Parallelism

SIMD on GPU

Trivial

Parallelism

An Intro to

MPI

Troop

Counting

Example

Line graph topology

Tree graph topology

Implementation

Parallel

Mandelbrot

Wrapping up

```
>> soldiers3
Worker 1:
    Soldier 1 reporting, sir! My neighbors are 2, sir!
Worker 2:
    Soldier 2 reporting, sir! My neighbors are 1, 3, 11, sir!
Worker 3:
    Soldier 3 reporting, sir! My neighbors are 2, 4, 5, sir!
Worker 4:
    Soldier 4 reporting, sir! My neighbors are 3, sir!
Worker 5:
    Soldier 5 reporting, sir! My neighbors are 3, sir!
Worker 6:
    Soldier 6 reporting, sir! My neighbors are 8, sir!
Worker 7:
    Soldier 7 reporting, sir! My neighbors are 8, sir!
Worker 8:
    Soldier 8 reporting, sir! My neighbors are 6, 7, 9, sir!
Worker 9:
    Soldier 9 reporting, sir! My neighbors are 8, 10, 12, sir!
Worker 10:
    Soldier 10 reporting, sir! My neighbors are 9, 11, sir!
Worker 11:
    Soldier 11 reporting, sir! My neighbors are 2, 10, sir!
Worker 12:
    Soldier 12 reporting, sir! My neighbors are 9, 13, 14, sir!
Worker 13:
    Soldier 13 reporting, sir! My neighbors are 12, sir!
Worker 14:
    Soldier 14 reporting, sir! My neighbors are 12, sir!
Worker 1:
    1 reached barrier.
```

# Sample output II

OS

Parallelism

SIMD on GPU

Trivial

Parallelism

An Intro to  
MPITroop  
Counting  
Example

Line graph topology

Tree graph topology

Implementation

Parallel  
Mandelbrot

Wrapping up

```
1 crossed barrier.
2 noticed not receiving from 1 sending 1 to 2...
1 received 13 from 2.
1 is done.
Worker 2:
2 reached barrier.
2 crossed barrier.
2 sees data available from 3...
2 received value 3 from 3.
2 sees data available from 1...
2 received value 1 from 1.
11 noticed not receiving from 2 sending 5 to 11...
2 received 9 from 11.
2 sending 13 to 1...
2 completed sending 13 to 1.
2 sending 11 to 3...
2 completed sending 11 to 3.
2 is done.
Worker 3:
3 reached barrier.
3 crossed barrier.
3 sees data available from 4...
3 received value 1 from 4.
3 sees data available from 5...
3 received value 1 from 5.
2 noticed not receiving from 3 sending 3 to 2...
3 received 11 from 2.
3 sending 13 to 4...
3 completed sending 13 to 4.
3 sending 13 to 5...
3 completed sending 13 to 5.
```

## Sample output III

OS

Parallelism

SIMD on GPU

Trivial

Parallelism

An Intro to

MPI

Troop

Counting

Example

Line graph topology

Tree graph topology

Implementation

Parallel

Mandelbrot

Wrapping up

```
3 is done.
Worker 4:
4 reached barrier.
4 crossed barrier.
3 noticed not receiving from 4 sending 1 to 3...
4 received 13 from 3.
4 is done.
Worker 5:
5 reached barrier.
5 crossed barrier.
3 noticed not receiving from 5 sending 1 to 3...
5 received 13 from 3.
5 is done.
Worker 6:
6 reached barrier.
6 crossed barrier.
8 noticed not receiving from 6 sending 1 to 8...
6 received 13 from 8.
6 is done.
Worker 7:
7 reached barrier.
7 crossed barrier.
8 noticed not receiving from 7 sending 1 to 8...
7 received 13 from 8.
7 is done.
Worker 8:
8 reached barrier.
8 crossed barrier.
8 sees data available from 7...
8 received value 1 from 7.
8 sees data available from 6...
```

## Sample output IV

OS

Parallelism

SIMD on GPU

Trivial

Parallelism

An Intro to

MPI

Troop

Counting

Example

Line graph topology

Tree graph topology

Implementation

Parallel

Mandelbrot

Wrapping up

```
      8 received value 1 from 6.
      9 noticed not receiving from 8 sending 3 to 9...
      8 received 11 from 9.
      8 sending 13 to 6...
      8 completed sending 13 to 6.
      8 sending 13 to 7...
      8 completed sending 13 to 7.
      8 is done.
Worker 9:
      9 reached barrier.
      9 crossed barrier.
      9 sees data available from 12...
      9 received value 3 from 12.
      9 sees data available from 8...
      9 received value 3 from 8.
      10 noticed not receiving from 9 sending 7 to 10...
      9 received 7 from 10.
      9 sending 11 to 8...
      9 completed sending 11 to 8.
      9 sending 11 to 12...
      9 completed sending 11 to 12.
COMMANDER 9 reporting count of 14.
      9 is done.
Worker 10:
      10 reached barrier.
      10 crossed barrier.
      10 sees data available from 9...
      10 received value 7 from 9.
      11 noticed not receiving from 10 sending 8 to 11...
      10 received 6 from 11.
      10 sending 7 to 9...
```

# Sample output V

OS

Parallelism

SIMD on GPU

Trivial

Parallelism

An Intro to  
MPITroop  
Counting  
Example

Line graph topology

Tree graph topology

Implementation

Parallel  
Mandelbrot

Wrapping up


```
10 completed sending 7 to 9.
10 is done.
Worker 11:
11 reached barrier.
11 crossed barrier.
11 sees data available from 2...
11 received value 5 from 2.
10 noticed not receiving from 11 sending 6 to 10...
11 received 8 from 10.
11 sending 9 to 2...
11 completed sending 9 to 2.
11 is done.
Worker 12:
12 reached barrier.
12 crossed barrier.
12 sees data available from 14...
12 received value 1 from 14.
12 sees data available from 13...
12 received value 1 from 13.
9 noticed not receiving from 12 sending 3 to 9...
12 received 11 from 9.
12 sending 13 to 13...
12 completed sending 13 to 13.
12 sending 13 to 14...
12 completed sending 13 to 14.
12 is done.
Worker 13:
13 reached barrier.
13 crossed barrier.
12 noticed not receiving from 13 sending 1 to 12...
13 received 13 from 12.
```

# Sample output VI

```
13 is done.  
Worker 14:  
14 reached barrier.  
14 crossed barrier.  
12 noticed not receiving from 14 sending 1 to 12...  
14 received 13 from 12.  
14 is done.  
Running total of 1 is 13.  
Running total of 2 is 13.  
Running total of 3 is 13.  
Running total of 4 is 13.  
Running total of 5 is 13.  
Running total of 6 is 13.  
Running total of 7 is 13.  
Running total of 8 is 13.  
Running total of 9 is 13.  
Running total of 10 is 13.  
Running total of 11 is 13.  
Running total of 12 is 13.  
Running total of 13 is 13.  
Running total of 14 is 13.  
>>
```

# Basic CPU Mandelbrot I

10.24secs (without GPU)





## Basic CPU Mandelbrot II

OS

Parallelism

SIMD on GPU

Trivial

Parallelism

An Intro to

MPI

Troop

Counting

Example

Line graph topology

Tree graph topology

Implementation

Parallel

Mandelbrot

Wrapping up

```

1  %-----
2  % File :      mandelbrot.m
3  %-----
4  %
5  % Author :    Marek Rychlik (rychlik@arizona.edu)
6  % Date :      Mon Feb 27 12:27:35 2023
7  % Copying :   (C) Marek Rychlik, 2020. All rights reserved.
8  %
9  %-----
10 % Mandelbrot without GPU or parallelization (MATLAB stock example)
11 maxIterations = 500;
12 gridSize = 1000;
13 xlim = [-0.748766713922161, -0.748766707771757];
14 ylim = [ 0.123640844894862,  0.123640851045266];
15
16 % Setup
17 t = tic();
18 x = linspace( xlim(1), xlim(2), gridSize );
19 y = linspace( ylim(1), ylim(2), gridSize );
20 [xGrid,yGrid] = meshgrid( x, y );
21 z0 = xGrid + 1i*yGrid;
22 count = ones( size(z0) );
23
24 % Calculate
25 z = z0;
26 for n = 0:maxIterations
27     z = z.*z + z0;
28     inside = abs( z ) <= 2;
29     count = count + inside;
30 end
31 count = log( count );

```

# Basic CPU Mandelbrot III

OS

Parallelism

SIMD on GPU

Trivial

Parallelism

An Intro to  
MPI

Troop

Counting

Example

Line graph topology

Tree graph topology

Implementation

Parallel  
Mandelbrot

Wrapping up


```
32
33 % Show
34 cpuTime = toc( t );
35 fig = gcf;
36 fig.Position = [200 200 600 600];
37 imagesc( x, y, count );
38 colormap( [jet();flipud( jet() );0 0 0] );
39 axis off
40 title( sprintf( '%1.2fsecs (without GPU)', cpuTime ) );
```

# Parallelized CPU Mandelbrot (8 workers) I

NOTE: Scaled up, higher resolution, larger radius  
(non-parallel time: 35 sec).

# Parallelized CPU Mandelbrot (8 workers) II

20.56 secs (without GPU, num. workers: 8, radius: 4.00, radius)



# Parallelized CPU Mandelbrot (8 workers) — code I

```
1 function count = mandel(x1,x2,y1,y2,gridSize,maxIterations,radius)
2     x = linspace( x1, x2, gridSize(1) );
3     y = linspace( y1, y2, gridSize(2) );
4     [xGrid,yGrid] = meshgrid( x, y );
5     z0 = xGrid + 1i*yGrid;
6     count = ones( size(z0) );
7
8     % Calculate
9     z = z0;
10    for n = 0:maxIterations
11        z = z.*z + z0;
12        inside = abs( z )<=radius;
13        count = count + inside;
14    end
15    count = log( count );
16 end
```

# Parallelized CPU Mandelbrot (8 workers) — code II

```

1  %-----
2  % File :      mandelbrotParallel.m
3  %-----
4  %
5  % Author :    Marek Rychlik (rychlik@arizona.edu)
6  % Date :      Mon Feb 27 12:32:56 2023
7  % Copying :   (C) Marek Rychlik, 2020. All rights reserved.
8  %
9  %-----
10 % Mandelbrot without GPU, parfor (MATLAB stock example, modified)
11 p=gcp('nocreate');
12 if isempty(p)
13     p = parpool('local', 8)
14 end
15 disp(sprintf('Number of workers: %d', p.NumWorkers));
16
17 maxIterations = 500;
18 gridSize = [2048,2048]; % Must be divisible by 8
19 radius=4;
20 xlim = [-0.748766713922161, -0.748766707771757];
21 ylim = [ 0.123640844894862, 0.123640851045266];
22
23 % Setup
24 x1 = xlim(1); x2=xlim(2); y1=ylim(1); y2=ylim(2);
25
26 % Non-parallel calculation
27 tic;count0 = mandel(x1, x2, y1, y2, gridSize, maxIterations, radius);disp('Non-
    parallel time');toc

```

# Parallelized CPU Mandelbrot (8 workers) — code III

```

28
29 numSlices = p.NumWorkers*32;
30 dx = (x2-x1) ./ numSlices;
31 gridSizeParallel = gridSize ./ [ numSlices , 1];
32 count = [];
33
34 q = Par(numSlices); % Par is a utility class for benchmarking
    parallel loops
35 parfor j=1:numSlices
36     Par.tic;
37     countLocal = mandel(x1 + (j-1).*dx, x1 + j.*dx, y1, y2, gridSizeParallel,
        maxIterations, radius);
38     count = [count, countLocal];
39     q(j)=Par.toc;
40 end
41 stop(q); plot(q);
42
43 % Show
44 cpuTime = q.StopTime;
45 fig = gcf;
46 fig.Position = [200 200 1024 1024];
47 imagesc( x, y, count );
48 colormap( [jet(); flipud( jet() ); 0 0 0] );
49 axis off
50 title( sprintf( '%1.2f secs (without GPU, num. workers: %d, radius: %1.2f,
    radius)', cpuTime, p.NumWorkers, radius ) );

```

# What has been left out?

- **OpenMP** (e.g., GOMP=GNU OpenMP); a high level interface to threads (SPMD) and vectorization (SIMD); realized as C/Fortran compiler **pragmas** (annotations)
- **POSIX threads** (“pthreads”); a C library available on most OS which allows direct access to multi-threading and thread synchronization
- Extensive **C++ language** constructs supporting parallelism
- **Building hardware**; hardware is inherently parallel; the most straightforward hardware to build is FPGA (Field-Programmable Gate Arrays); programming languages **Verilog** and **VHDL**
- University of Arizona **HPC facilities**